

Part 1 Question 2: Analyze the avocado data

Tasks:

```
In [ ]: import pandas as pd
```

1. Read the data from the CSV file into a DataFrame.

```
In [ ]: df = pd.read_csv('datasets/avocado.csv')
df
```

```
Out [ ]:
```

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total B...
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014

18249 rows × 14 columns

2. Display type memory consumption and null count information using the info() method.

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0             18249 non-null  int64
1   Date                   18249 non-null  object
2   AveragePrice           18249 non-null  float64
3   Total Volume           18249 non-null  float64
4   4046                   18249 non-null  float64
5   4225                   18249 non-null  float64
6   4770                   18249 non-null  float64
7   Total Bags             18249 non-null  float64
8   Small Bags             18249 non-null  float64
9   Large Bags             18249 non-null  float64
10  XLarge Bags            18249 non-null  float64
11  type                   18249 non-null  object
12  year                   18249 non-null  int64
13  region                 18249 non-null  object
dtypes: float64(9), int64(2), object(3)
memory usage: 1.9+ MB
```

3. Display the number of unique values in each column.

```
In [ ]: df.nunique()
```

```
Out[ ]: Unnamed: 0      53
Date                169
AveragePrice        259
Total Volume        18237
4046                17702
4225                18103
4770                12071
Total Bags          18097
Small Bags          17321
Large Bags          15082
XLarge Bags         5588
type                2
year                4
region              54
dtype: int64
```

4. Display all the rows of data that JupyterLab displays by default.

```
In [ ]: df
```

Out []:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Tc Bz
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145
3	3	2015-12-06	1.08	78992.15	1132.00	71976.41	72.58	5811
4	4	2015-11-29	1.28	51039.60	941.48	43838.39	75.78	6183
...
18244	7	2018-02-04	1.63	17074.83	2046.96	1529.20	0.00	13498
18245	8	2018-01-28	1.71	13888.04	1191.70	3431.50	0.00	9264
18246	9	2018-01-21	1.87	13766.76	1191.92	2452.79	727.94	9394
18247	10	2018-01-14	1.93	16205.22	1527.63	2981.04	727.01	10969
18248	11	2018-01-07	1.62	17489.58	2894.77	2356.13	224.53	12014

18249 rows x 14 columns

5. Display the first and last five rows of data and the first and last four columns of data.

```
In [ ]: pd.set_option('display.max_rows', 10)
pd.set_option('display.max_columns', 8)
df
```

Out []:

	Unnamed: 0	Date	AveragePrice	Total Volume	...	XLarge Bags	type	year
0	0	2015-12-27	1.33	64236.62	...	0.0	conventional	2015
1	1	2015-12-20	1.35	54876.98	...	0.0	conventional	2015
2	2	2015-12-13	0.93	118220.22	...	0.0	conventional	2015
3	3	2015-12-06	1.08	78992.15	...	0.0	conventional	2015
4	4	2015-11-29	1.28	51039.60	...	0.0	conventional	2015
...
18244	7	2018-02-04	1.63	17074.83	...	0.0	organic	2018 We
18245	8	2018-01-28	1.71	13888.04	...	0.0	organic	2018 We
18246	9	2018-01-21	1.87	13766.76	...	0.0	organic	2018 We
18247	10	2018-01-14	1.93	16205.22	...	0.0	organic	2018 We
18248	11	2018-01-07	1.62	17489.58	...	0.0	organic	2018 We

18249 rows x 14 columns

6. Choose any three columns access them with bracket notation and display the first five rows of this data.

In []: `df[['Date', 'AveragePrice', 'Total Volume']].head(5)`

Out []:

	Date	AveragePrice	Total Volume
0	2015-12-27	1.33	64236.62
1	2015-12-20	1.35	54876.98
2	2015-12-13	0.93	118220.22
3	2015-12-06	1.08	78992.15
4	2015-11-29	1.28	51039.60

7. Select one column and access it with dot notation.

```
In [ ]: df.Date
```

```
Out [ ]: 0      2015-12-27
         1      2015-12-20
         2      2015-12-13
         3      2015-12-06
         4      2015-11-29
         ...
        18244    2018-02-04
        18245    2018-01-28
        18246    2018-01-21
        18247    2018-01-14
        18248    2018-01-07
        Name: Date, Length: 18249, dtype: object
```

8. Multiply the Total Volume and AveragePrice columns and store the result in a new column called EstimatedRevenue. Then display the first five rows of this data to confirm that the column was added and has the correct values.

```
In [ ]: df['EstimatedRevenue'] = df['AveragePrice'] * df['Total Volume']
        df.head(5)
```

```
Out [ ]:   Unnamed: 0  Date  AveragePrice  Total Volume  ...  type  year  region  Estimate
```

	Unnamed: 0	Date	AveragePrice	Total Volume	...	type	year	region	Estimate
0	0	2015-12-27	1.33	64236.62	...	conventional	2015	Albany	8
1	1	2015-12-20	1.35	54876.98	...	conventional	2015	Albany	7
2	2	2015-12-13	0.93	118220.22	...	conventional	2015	Albany	109
3	3	2015-12-06	1.08	78992.15	...	conventional	2015	Albany	8
4	4	2015-11-29	1.28	51039.60	...	conventional	2015	Albany	6

5 rows x 15 columns

9. Create a DataFrame that's grouped by region and type and that includes the average price for the grouped columns. Then reset the index and display the first five rows.

```
In [ ]: df[['region', 'type', 'AveragePrice']].groupby(['region', 'type']).mean().re
```


Out []:

	region	type	AveragePrice
0	Albany	conventional	1.348757
1	Albany	organic	1.773314
2	Atlanta	conventional	1.068817
3	Atlanta	organic	1.607101
4	BaltimoreWashington	conventional	1.344201

10. Create a bar plot that shows the mean median and standard deviation of the Total Volume column by year.

```
In [ ]: df['Year'] = pd.to_datetime(df['Date']).dt.year
plot_df = df[['Year', 'AveragePrice']].groupby('Year').agg(['mean', 'median', 'std'])
plot_df.plot.bar(x='Year', y='AveragePrice')
```

Out []: <Axes: xlabel='Year'>



Questions:

1. How many unique regions are there?

```
In [ ]: df['region'].nunique()
```

```
Out [ ]: 54
```

There are 54 unique regions in the dataset

2. What is the average price for each type of avocado (organic and conventional)? Be sure to include just the type and AveragePrice columns in the results.

```
In [ ]: df[['type', 'AveragePrice']].groupby('type').mean().reset_index()
```

```
Out [ ]:
```

	type	AveragePrice
0	conventional	1.158040
1	organic	1.653999

The average price for each type of avocado is as follows:

1. conventional - \$1.16
2. organic - \$1.65

3. Which region has the lowest average price for organic avocados? Hint: Create wide data from the grouped data that you created in task 8.

```
In [ ]: wide_df = df[['region', 'type', 'AveragePrice']].groupby(['region', 'type'])
wide_df[wide_df['type'] == 'organic'].sort_values('AveragePrice', ascending=
```

```
Out [ ]:
```

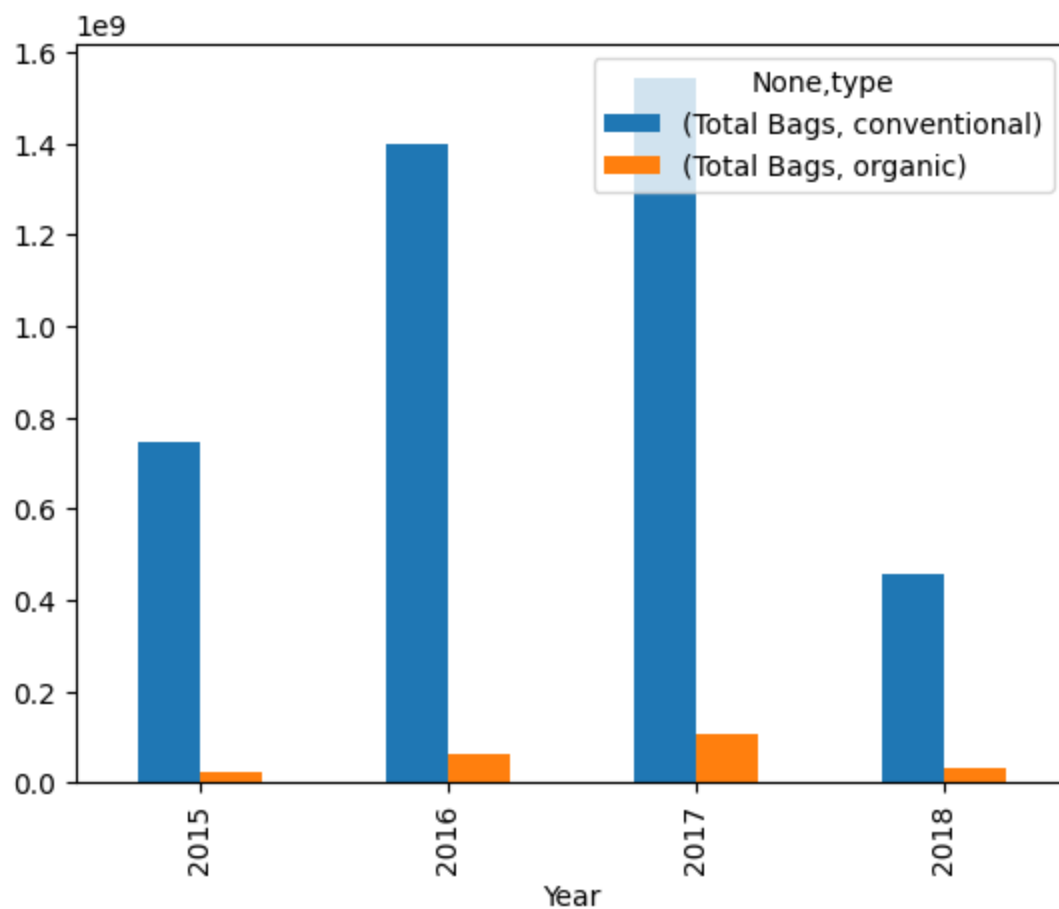
	region	type	AveragePrice
37	Houston	organic	1.270769

The region with the lowest average organic avocado price is Houston at \$1.27

4. Have the Total Bags sold per year of each type of avocado become more or less consistent over time?

```
In [ ]: plot_df = df[['Year', 'type', 'Total Bags']].groupby(['Year', 'type']).sum()
plot_df.plot.bar()
```

```
Out [ ]: <Axes: xlabel='Year'>
```



The graph above represents the total bags sold per year of each type of avocado. From this graph it is clear that the total bags sold per year has become less consistent over time, this can be seen with the significant drop from 2017 to 2018.