

# Assignment 4 - Multiclass Classifier Neural Network

```
In [ ]: import pandas as pd
```

## Step 1: Load the Data

```
In [ ]: df = pd.read_csv('a4-data/train.csv')
df.head()
```

```
Out [ ]:
```

	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5	Feature_6	Feature_7	Fea
0	3289	22	19	240	93	1708	205	
1	2963	21	18	134	27	1243	206	
2	3037	185	9	127	10	6462	222	
3	3113	203	13	190	22	2125	213	
4	3128	346	9	120	36	552	203	

5 rows x 55 columns

we need to check for any null or missing values and deal with them accordingly, here we have none

```
In [ ]: df.isnull().sum().sum()
```

```
Out [ ]: 0
```

## Step 2: Split the Data into Features (X) and Target (Y)

Here we split the training data into X and y variables where X contains a feature matrix and y contains a target vector, this is done so the model has a target value to compare to its predictions, this enables us to calculate the loss which ultimately allows the model to learn.

```
In [ ]: # split dataset into features and target
X = df.drop('Target', axis=1)
y = df['Target']
X.shape, y.shape
```

```
Out [ ]: ((464809, 54), (464809,))
```

Now that we have our target and features separated, we need to normalize the feature matrix which will help the model converge to a minima faster by leveling out the gradient. The reason for not normalizing the target vector is because it contains categorical data which should not be normalized.

## Normalize the features

When no normalization is done the models accuracy is poor with a value less than 60%, but when all features are normalized the models performance dramatically increases to over 90%

We can take this one step further by only normalizing the non-binary features which improves the models performance even more, even when the model has a significantly higher accuracy of 93.6%, only normalizing the non-binary features gives us a 1% increase with a 94.6%

Here we get only the non-binary columns and normalize them using z-score normalization

```
In [ ]: binary_columns = [col for col in df.columns if df[col].dropna().isin([0, 1])
non_binary_columns = [col for col in X.columns if col not in binary_columns]

# Normalize non-binary columns
X[non_binary_columns] = (X[non_binary_columns] - X[non_binary_columns].mean(
X.head())
```

```
Out [ ]: e_10  ...  Feature_45  Feature_46  Feature_47  Feature_48  Feature_49  Feature_50  F
5718  ...      0      0      1      0      0      0
1309  ...      0      0      0      0      0      0
3933  ...      0      0      0      0      0      0
3645  ...      0      1      0      0      0      0
7276  ...      0      0      0      0      0      0
```

## Step 3: Data Splitting

We need to split our data into training and test datasets so that we can evaluate and validate our models performance, if we tested on the entire dataset, we would not be testing the model on unseen data which would not be a valid test of its performance.

```
In [ ]: from sklearn.model_selection import train_test_split

# Split dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

Out[ ]: ((325366, 54), (139443, 54), (325366,), (139443,))
```

## Step 4: Build the Neural Network Model

Here we build the model. In order to have optimal parameters, I decided to use the Hyperband algorithm to tune my parameters using the kerastuner library. I first started with a small depth and width, letting the algorithm find the optimal learning rate. The smaller models performed well, but by having a larger neural network lead to better models.

The best model from the Hyperband process has the following parameters:

- 54 node input layer, for the 54 features in X
- 120 node hidden layer, relu activation
- 112 node hidden layer, relu activation
- 104 node hidden layer, relu activation
- 80 node hidden layer, relu activation
- 7 node output layer, softmax activation
- learned rate of 0.00085410586

I tested using few layers but ultimately ended up with 4 hidden layers which seemed to give good results for the model

Relu was used as the activation function for hidden layers because it introduces non-linearity into the model which allows the model to learn complex patterns and relationships which is what we are trying to achieve with our classifier.

Softmax was used for the output layers activation function as it converts the raw output from a neural network into probabilities which we can then select the class with the highest probability, this is perfect for our classifier's output layer.

Since the number of nodes in the output layer of the model needs to match the range for the target, we use the target's unique value count plus 1 as the output layer's shape

Here we set up the model with our found widths for each layer from the hyperparameter tuning

```
In [ ]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Input
```

```
# Define the model
model = Sequential([
    Input(shape=(X_train.shape[1],)),
    Dense(units=120, activation='relu'),
    Dense(units=112, activation='relu'),
    Dense(units=104, activation='relu'),
    Dense(units=80, activation='relu'),
    Dense(units=(y_train.nunique()+1), activation='softmax')
])
```

```
/Users/bradymitchelmore/Library/Mobile Documents/com~apple~CloudDocs/MUN/Year 3/Term 7/COMP 3401/Assignments/comp3401/lib/python3.9/site-packages/urllib3/__init__.py:35: NotOpenSSLWarning: urllib3 v2 only supports OpenSSL 1.1.1+, currently the 'ssl' module is compiled with 'LibreSSL 2.8.3'. See: https://github.com/urllib3/urllib3/issues/3020
warnings.warn(
```

## Step 5: Train the Model

Here we train the model using the learning rate from our hyperparameter tuning and 100 epochs


- I choose the Adam optimizer due to its adaptive learning rates and efficiency which help the model converge faster.
- The loss function used was `sparse_categorical_crossentropy`, which is a function that is good when dealing with multiclass classification where each sample belongs to one of many classes and the labels are numeric, these are 2 properties of our target vector which deems this loss function suitable.


```
In [ ]: from tensorflow.keras.optimizers import Adam
        from tensorflow.keras.optimizers.schedules import ExponentialDecay


        initial_learning_rate = 0.00085410586


        optimizer = Adam(learning_rate=initial_learning_rate)


        model.compile(loss='sparse_categorical_crossentropy', optimizer=optimizer, metrics=['accuracy'])
        history = model.fit(X_train, y_train, epochs=100, validation_data=(X_test, y_test))
```


Epoch 1/100  
**10168/10168**  **9s** 828us/step - accuracy: 0.7519 - loss: 0.5861 - val\_accuracy: 0.8319 - val\_loss: 0.4034


Epoch 2/100  
**10168/10168**  **8s** 761us/step - accuracy: 0.8403 - loss: 0.3805 - val\_accuracy: 0.8606 - val\_loss: 0.3378


Epoch 3/100  
**10168/10168**  **7s** 702us/step - accuracy: 0.8671 - loss: 0.3204 - val\_accuracy: 0.8728 - val\_loss: 0.3068


Epoch 4/100  
**10168/10168**  **8s** 798us/step - accuracy: 0.8812 - loss: 0.2870 - val\_accuracy: 0.8860 - val\_loss: 0.2772


Epoch 5/100  
**10168/10168**  **7s** 733us/step - accuracy: 0.8916 - loss: 0.2620 - val\_accuracy: 0.8945 - val\_loss: 0.2582


Epoch 6/100  
**10168/10168**  **7s** 714us/step - accuracy: 0.8991 - loss: 0.2450 - val\_accuracy: 0.8917 - val\_loss: 0.2592


Epoch 7/100  
**10168/10168**  **7s** 694us/step - accuracy: 0.9059 - loss: 0.2303 - val\_accuracy: 0.9028 - val\_loss: 0.2371


Epoch 8/100  
**10168/10168**  **7s** 720us/step - accuracy: 0.9100 - loss: 0.2191 - val\_accuracy: 0.9088 - val\_loss: 0.2243


Epoch 9/100  
**10168/10168**  **7s** 705us/step - accuracy: 0.9143 - loss: 0.2079 - val\_accuracy: 0.9135 - val\_loss: 0.2169


Epoch 10/100  
**10168/10168**  **7s** 703us/step - accuracy: 0.9180 - loss: 0.2021 - val\_accuracy: 0.9110 - val\_loss: 0.2208


Epoch 11/100  
**10168/10168**  **7s** 707us/step - accuracy: 0.9200 - loss: 0.1955 - val\_accuracy: 0.9129 - val\_loss: 0.2159


Epoch 12/100  
**10168/10168**  **8s** 738us/step - accuracy: 0.9237 - loss: 0.1872 - val\_accuracy: 0.9168 - val\_loss: 0.2066


Epoch 13/100  
**10168/10168**  **7s** 697us/step - accuracy: 0.9260 - loss: 0.1824 - val\_accuracy: 0.9192 - val\_loss: 0.2006


Epoch 14/100  
**10168/10168**  **7s** 722us/step - accuracy: 0.9273 - loss: 0.1787 - val\_accuracy: 0.9218 - val\_loss: 0.1959
















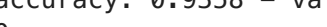
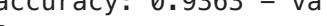
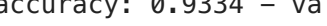
Epoch 15/100  
**10168/10168**  **7s** 721us/step - accuracy: 0.9293 - loss: 0.1737 - val\_accuracy: 0.9226 - val\_loss: 0.1965


Epoch 16/100  
**10168/10168**  **7s** 696us/step - accuracy: 0.9310 - loss: 0.1699 - val\_accuracy: 0.9198 - val\_loss: 0.2021


Epoch 17/100  
**10168/10168**  **8s** 740us/step - accuracy: 0.9315 - loss: 0.1684 - val\_accuracy: 0.9244 - val\_loss: 0.1916


Epoch 18/100  
**10168/10168**  **7s** 688us/step - accuracy: 0.9341 - loss: 0.1641 - val\_accuracy: 0.9255 - val\_loss: 0.1894


Epoch 19/100  
**10168/10168**  **7s** 727us/step - accuracy: 0.9350 - loss: 0.


1611 - val\_accuracy: 0.9241 - val\_loss: 0.1950  
Epoch 20/100  
**10168/10168**  7s 715us/step - accuracy: 0.9351 - loss: 0.  
1596 - val\_accuracy: 0.9276 - val\_loss: 0.1831  
Epoch 21/100  
**10168/10168**  7s 693us/step - accuracy: 0.9369 - loss: 0.  
1557 - val\_accuracy: 0.9286 - val\_loss: 0.1809  
Epoch 22/100  
**10168/10168**  7s 717us/step - accuracy: 0.9375 - loss: 0.  
1547 - val\_accuracy: 0.9288 - val\_loss: 0.1833  
Epoch 23/100  
**10168/10168**  7s 717us/step - accuracy: 0.9384 - loss: 0.  
1524 - val\_accuracy: 0.9293 - val\_loss: 0.1805  
Epoch 24/100  
**10168/10168**  7s 703us/step - accuracy: 0.9394 - loss: 0.  
1499 - val\_accuracy: 0.9314 - val\_loss: 0.1778  
Epoch 25/100  
**10168/10168**  7s 704us/step - accuracy: 0.9403 - loss: 0.  
1472 - val\_accuracy: 0.9307 - val\_loss: 0.1813  
Epoch 26/100  
**10168/10168**  7s 696us/step - accuracy: 0.9410 - loss: 0.  
1463 - val\_accuracy: 0.9315 - val\_loss: 0.1772  
Epoch 27/100  
**10168/10168**  7s 718us/step - accuracy: 0.9420 - loss: 0.  
1423 - val\_accuracy: 0.9343 - val\_loss: 0.1699  
Epoch 28/100  
**10168/10168**  7s 724us/step - accuracy: 0.9431 - loss: 0.  
1426 - val\_accuracy: 0.9317 - val\_loss: 0.1796  
Epoch 29/100  
**10168/10168**  7s 685us/step - accuracy: 0.9434 - loss: 0.  
1408 - val\_accuracy: 0.9335 - val\_loss: 0.1719  
Epoch 30/100  
**10168/10168**  7s 716us/step - accuracy: 0.9430 - loss: 0.  
1398 - val\_accuracy: 0.9333 - val\_loss: 0.1726  
Epoch 31/100  
**10168/10168**  7s 713us/step - accuracy: 0.9442 - loss: 0.  
1384 - val\_accuracy: 0.9308 - val\_loss: 0.1815  
Epoch 32/100  
**10168/10168**  8s 746us/step - accuracy: 0.9441 - loss: 0.  
1378 - val\_accuracy: 0.9349 - val\_loss: 0.1718  
Epoch 33/100  
**10168/10168**  8s 746us/step - accuracy: 0.9457 - loss: 0.  
1343 - val\_accuracy: 0.9318 - val\_loss: 0.1847  
Epoch 34/100  
**10168/10168**  7s 729us/step - accuracy: 0.9460 - loss: 0.  
1350 - val\_accuracy: 0.9358 - val\_loss: 0.1668  
Epoch 35/100  
**10168/10168**  8s 754us/step - accuracy: 0.9462 - loss: 0.  
1335 - val\_accuracy: 0.9363 - val\_loss: 0.1707  
Epoch 36/100  
**10168/10168**  8s 766us/step - accuracy: 0.9469 - loss: 0.  
1313 - val\_accuracy: 0.9334 - val\_loss: 0.1779  
Epoch 37/100  
**10168/10168**  7s 717us/step - accuracy: 0.9476 - loss: 0.  
1301 - val\_accuracy: 0.9369 - val\_loss: 0.1654  
Epoch 38/100


**10168/10168**  **7s** 728us/step - accuracy: 0.9482 - loss: 0.1291 - val\_accuracy: 0.9380 - val\_loss: 0.1677  
Epoch 39/100


**10168/10168**  **7s** 729us/step - accuracy: 0.9476 - loss: 0.1295 - val\_accuracy: 0.9374 - val\_loss: 0.1662  
Epoch 40/100


**10168/10168**  **7s** 727us/step - accuracy: 0.9486 - loss: 0.1284 - val\_accuracy: 0.9372 - val\_loss: 0.1668  
Epoch 41/100


**10168/10168**  **7s** 714us/step - accuracy: 0.9487 - loss: 0.1282 - val\_accuracy: 0.9371 - val\_loss: 0.1678  
Epoch 42/100


**10168/10168**  **7s** 718us/step - accuracy: 0.9498 - loss: 0.1253 - val\_accuracy: 0.9390 - val\_loss: 0.1633  
Epoch 43/100


**10168/10168**  **7s** 722us/step - accuracy: 0.9494 - loss: 0.1260 - val\_accuracy: 0.9373 - val\_loss: 0.1652  
Epoch 44/100


**10168/10168**  **8s** 737us/step - accuracy: 0.9509 - loss: 0.1236 - val\_accuracy: 0.9386 - val\_loss: 0.1683  
Epoch 45/100


**10168/10168**  **7s** 729us/step - accuracy: 0.9504 - loss: 0.1242 - val\_accuracy: 0.9365 - val\_loss: 0.1696  
Epoch 46/100


**10168/10168**  **8s** 750us/step - accuracy: 0.9503 - loss: 0.1231 - val\_accuracy: 0.9393 - val\_loss: 0.1655  
Epoch 47/100


**10168/10168**  **7s** 718us/step - accuracy: 0.9506 - loss: 0.1234 - val\_accuracy: 0.9389 - val\_loss: 0.1628  
Epoch 48/100


**10168/10168**  **7s** 721us/step - accuracy: 0.9516 - loss: 0.1206 - val\_accuracy: 0.9372 - val\_loss: 0.1709  
Epoch 49/100


**10168/10168**  **7s** 719us/step - accuracy: 0.9517 - loss: 0.1207 - val\_accuracy: 0.9361 - val\_loss: 0.1751  
Epoch 50/100


**10168/10168**  **8s** 739us/step - accuracy: 0.9520 - loss: 0.1200 - val\_accuracy: 0.9402 - val\_loss: 0.1628  
Epoch 51/100

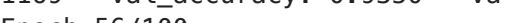
**10168/10168**  **7s** 728us/step - accuracy: 0.9528 - loss: 0.1183 - val\_accuracy: 0.9405 - val\_loss: 0.1608  
Epoch 52/100




















**10168/10168**  **8s** 737us/step - accuracy: 0.9523 - loss: 0.1201 - val\_accuracy: 0.9401 - val\_loss: 0.1643  
Epoch 53/100

**10168/10168**  **7s** 706us/step - accuracy: 0.9529 - loss: 0.1188 - val\_accuracy: 0.9393 - val\_loss: 0.1652  
Epoch 54/100



















**10168/10168**  **7s** 732us/step - accuracy: 0.9529 - loss: 0.1189 - val\_accuracy: 0.9399 - val\_loss: 0.1639  
Epoch 55/100

**10168/10168**  **8s** 765us/step - accuracy: 0.9533 - loss: 0.1169 - val\_accuracy: 0.9350 - val\_loss: 0.1754  
Epoch 56/100

**10168/10168**  **8s** 738us/step - accuracy: 0.9535 - loss: 0.1172 - val\_accuracy: 0.9405 - val\_loss: 0.1600

Epoch 57/100  
**10168/10168**  **7s** 723us/step - accuracy: 0.9539 - loss: 0.1160 - val\_accuracy: 0.9374 - val\_loss: 0.1715  
Epoch 58/100  
**10168/10168**  **8s** 736us/step - accuracy: 0.9540 - loss: 0.1162 - val\_accuracy: 0.9399 - val\_loss: 0.1669  
Epoch 59/100  
**10168/10168**  **8s** 763us/step - accuracy: 0.9547 - loss: 0.1138 - val\_accuracy: 0.9381 - val\_loss: 0.1702  
Epoch 60/100  
**10168/10168**  **8s** 792us/step - accuracy: 0.9541 - loss: 0.1155 - val\_accuracy: 0.9413 - val\_loss: 0.1637  
Epoch 61/100  
**10168/10168**  **7s** 723us/step - accuracy: 0.9542 - loss: 0.1149 - val\_accuracy: 0.9375 - val\_loss: 0.1742  
Epoch 62/100  
**10168/10168**  **8s** 769us/step - accuracy: 0.9545 - loss: 0.1149 - val\_accuracy: 0.9376 - val\_loss: 0.1782  
Epoch 63/100  
**10168/10168**  **8s** 749us/step - accuracy: 0.9546 - loss: 0.1138 - val\_accuracy: 0.9389 - val\_loss: 0.1671  
Epoch 64/100  
**10168/10168**  **7s** 734us/step - accuracy: 0.9550 - loss: 0.1137 - val\_accuracy: 0.9395 - val\_loss: 0.1715  
Epoch 65/100  
**10168/10168**  **7s** 701us/step - accuracy: 0.9541 - loss: 0.1153 - val\_accuracy: 0.9416 - val\_loss: 0.1642  
Epoch 66/100  
**10168/10168**  **8s** 773us/step - accuracy: 0.9552 - loss: 0.1129 - val\_accuracy: 0.9399 - val\_loss: 0.1694  
Epoch 67/100  
**10168/10168**  **8s** 737us/step - accuracy: 0.9554 - loss: 0.1129 - val\_accuracy: 0.9391 - val\_loss: 0.1754  
Epoch 68/100  
**10168/10168**  **8s** 738us/step - accuracy: 0.9552 - loss: 0.1137 - val\_accuracy: 0.9394 - val\_loss: 0.1702  
Epoch 69/100  
**10168/10168**  **7s** 726us/step - accuracy: 0.9556 - loss: 0.1114 - val\_accuracy: 0.9435 - val\_loss: 0.1588  
Epoch 70/100  
**10168/10168**  **9s** 874us/step - accuracy: 0.9555 - loss: 0.1113 - val\_accuracy: 0.9384 - val\_loss: 0.1705  
Epoch 71/100  
**10168/10168**  **7s** 730us/step - accuracy: 0.9558 - loss: 0.1117 - val\_accuracy: 0.9401 - val\_loss: 0.1722  
Epoch 72/100  
**10168/10168**  **8s** 755us/step - accuracy: 0.9565 - loss: 0.1108 - val\_accuracy: 0.9415 - val\_loss: 0.1645  
Epoch 73/100  
**10168/10168**  **8s** 769us/step - accuracy: 0.9563 - loss: 0.1107 - val\_accuracy: 0.9414 - val\_loss: 0.1677  
Epoch 74/100  
**10168/10168**  **8s** 736us/step - accuracy: 0.9565 - loss: 0.1094 - val\_accuracy: 0.9413 - val\_loss: 0.1753  
Epoch 75/100  
**10168/10168**  **7s** 719us/step - accuracy: 0.9565 - loss: 0.



1109 - val\_accuracy: 0.9426 - val\_loss: 0.1622  
Epoch 76/100  
**10168/10168**  8s 788us/step - accuracy: 0.9570 - loss: 0.  
1100 - val\_accuracy: 0.9422 - val\_loss: 0.1637  
Epoch 77/100  
**10168/10168**  8s 752us/step - accuracy: 0.9568 - loss: 0.  
1097 - val\_accuracy: 0.9416 - val\_loss: 0.1683  
Epoch 78/100  
**10168/10168**  8s 740us/step - accuracy: 0.9561 - loss: 0.  
1122 - val\_accuracy: 0.9422 - val\_loss: 0.1663  
Epoch 79/100  
**10168/10168**  7s 730us/step - accuracy: 0.9580 - loss: 0.  
1079 - val\_accuracy: 0.9435 - val\_loss: 0.1624  
Epoch 80/100  
**10168/10168**  8s 778us/step - accuracy: 0.9577 - loss: 0.  
1071 - val\_accuracy: 0.9432 - val\_loss: 0.1643  
Epoch 81/100  
**10168/10168**  7s 734us/step - accuracy: 0.9566 - loss: 0.  
1109 - val\_accuracy: 0.9418 - val\_loss: 0.1727  
Epoch 82/100  
**10168/10168**  7s 734us/step - accuracy: 0.9576 - loss: 0.  
1069 - val\_accuracy: 0.9437 - val\_loss: 0.1625  
Epoch 83/100  
**10168/10168**  8s 737us/step - accuracy: 0.9567 - loss: 0.  
1119 - val\_accuracy: 0.9442 - val\_loss: 0.1637  
Epoch 84/100  
**10168/10168**  8s 737us/step - accuracy: 0.9579 - loss: 0.  
1081 - val\_accuracy: 0.9419 - val\_loss: 0.1675  
Epoch 85/100  
**10168/10168**  7s 729us/step - accuracy: 0.9575 - loss: 0.  
1077 - val\_accuracy: 0.9405 - val\_loss: 0.1723  
Epoch 86/100  
**10168/10168**  8s 750us/step - accuracy: 0.9578 - loss: 0.  
1083 - val\_accuracy: 0.9436 - val\_loss: 0.1630  
Epoch 87/100  
**10168/10168**  7s 736us/step - accuracy: 0.9578 - loss: 0.  
1079 - val\_accuracy: 0.9408 - val\_loss: 0.1696  
Epoch 88/100  
**10168/10168**  8s 742us/step - accuracy: 0.9584 - loss: 0.  
1073 - val\_accuracy: 0.9429 - val\_loss: 0.1687  
Epoch 89/100  
**10168/10168**  8s 736us/step - accuracy: 0.9588 - loss: 0.  
1061 - val\_accuracy: 0.9431 - val\_loss: 0.1680  
Epoch 90/100  
**10168/10168**  7s 732us/step - accuracy: 0.9588 - loss: 0.  
1056 - val\_accuracy: 0.9432 - val\_loss: 0.1633  
Epoch 91/100  
**10168/10168**  7s 727us/step - accuracy: 0.9575 - loss: 0.  
1093 - val\_accuracy: 0.9419 - val\_loss: 0.1749  
Epoch 92/100  
**10168/10168**  7s 732us/step - accuracy: 0.9587 - loss: 0.  
1055 - val\_accuracy: 0.9447 - val\_loss: 0.1627  
Epoch 93/100  
**10168/10168**  8s 744us/step - accuracy: 0.9590 - loss: 0.  
1051 - val\_accuracy: 0.9434 - val\_loss: 0.1654  
Epoch 94/100

```

10168/10168 ————— 8s 753us/step – accuracy: 0.9584 – loss: 0.
1081 – val_accuracy: 0.9427 – val_loss: 0.1679
Epoch 95/100
10168/10168 ————— 7s 736us/step – accuracy: 0.9578 – loss: 0.
1091 – val_accuracy: 0.9406 – val_loss: 0.1759
Epoch 96/100
10168/10168 ————— 8s 741us/step – accuracy: 0.9588 – loss: 0.
1056 – val_accuracy: 0.9408 – val_loss: 0.1734
Epoch 97/100
10168/10168 ————— 7s 734us/step – accuracy: 0.9583 – loss: 0.
1057 – val_accuracy: 0.9435 – val_loss: 0.1681
Epoch 98/100
10168/10168 ————— 8s 736us/step – accuracy: 0.9584 – loss: 0.
1063 – val_accuracy: 0.9418 – val_loss: 0.1713
Epoch 99/100
10168/10168 ————— 8s 747us/step – accuracy: 0.9583 – loss: 0.
1071 – val_accuracy: 0.9435 – val_loss: 0.1696
Epoch 100/100
10168/10168 ————— 8s 806us/step – accuracy: 0.9595 – loss: 0.
1051 – val_accuracy: 0.9444 – val_loss: 0.1742

```

## Step 6: Model Evaluation

Evaluating the model on the test set allows us to assess its performance on unseen data, giving us an indication of its generalization ability. Accuracy is used as a metric to quantify the percentage of correctly predicted instances, while the loss metric gives us an indication of how close our predictions are.

```
In [ ]: loss, accuracy = model.evaluate(X_test, y_test)
        print(f"Loss: {loss}, Accuracy: {accuracy}")
```

```

4358/4358 ————— 2s 338us/step – accuracy: 0.9436 – loss: 0.17
69
Loss: 0.17416204512119293, Accuracy: 0.9443858861923218

```

Here we can see our model gives us a 94.36% accuracy on the test set with a loss of just 0.1705

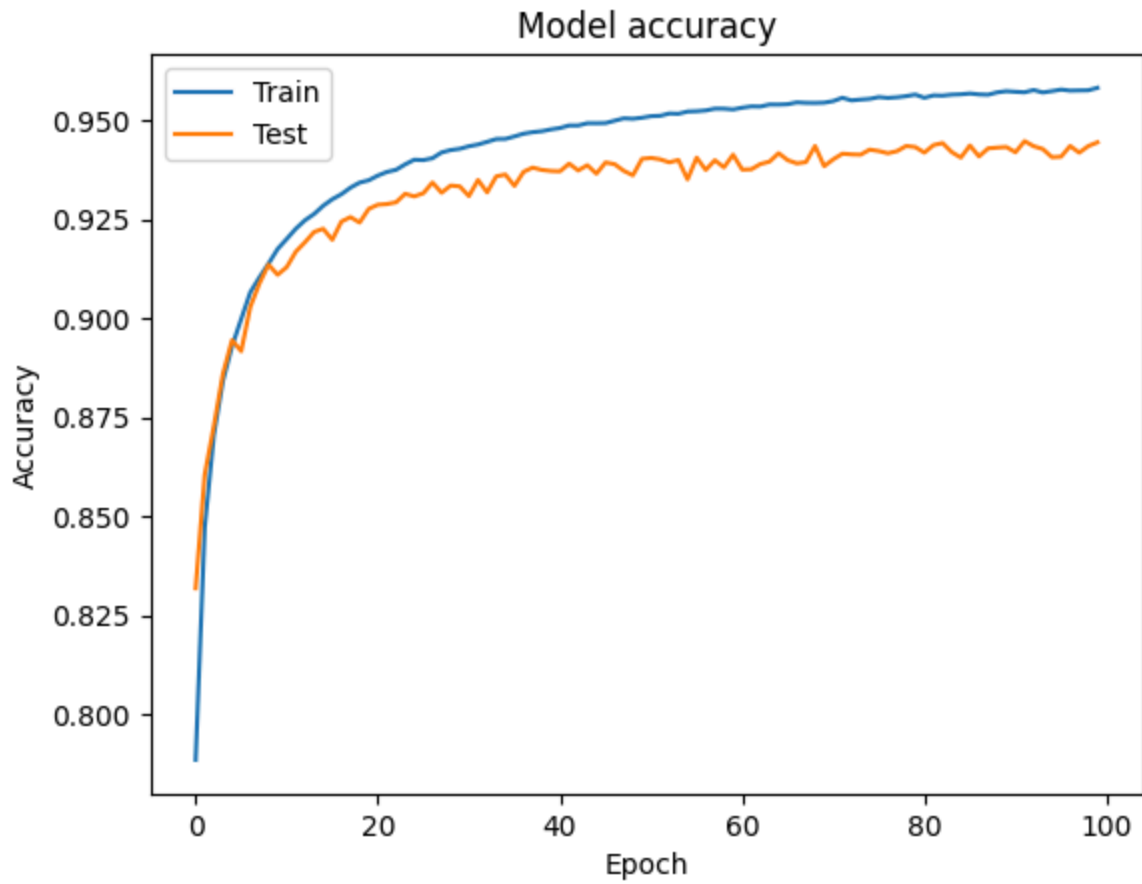
When training we used `X_test` and `y_test` as a validation set which gave us the `val_accuracy` and `val_loss` which measure the accuracy and loss in terms of the test set. We can plot the validation vs the training accuracy and loss to visualize how the model fits to its training data better than unseen data

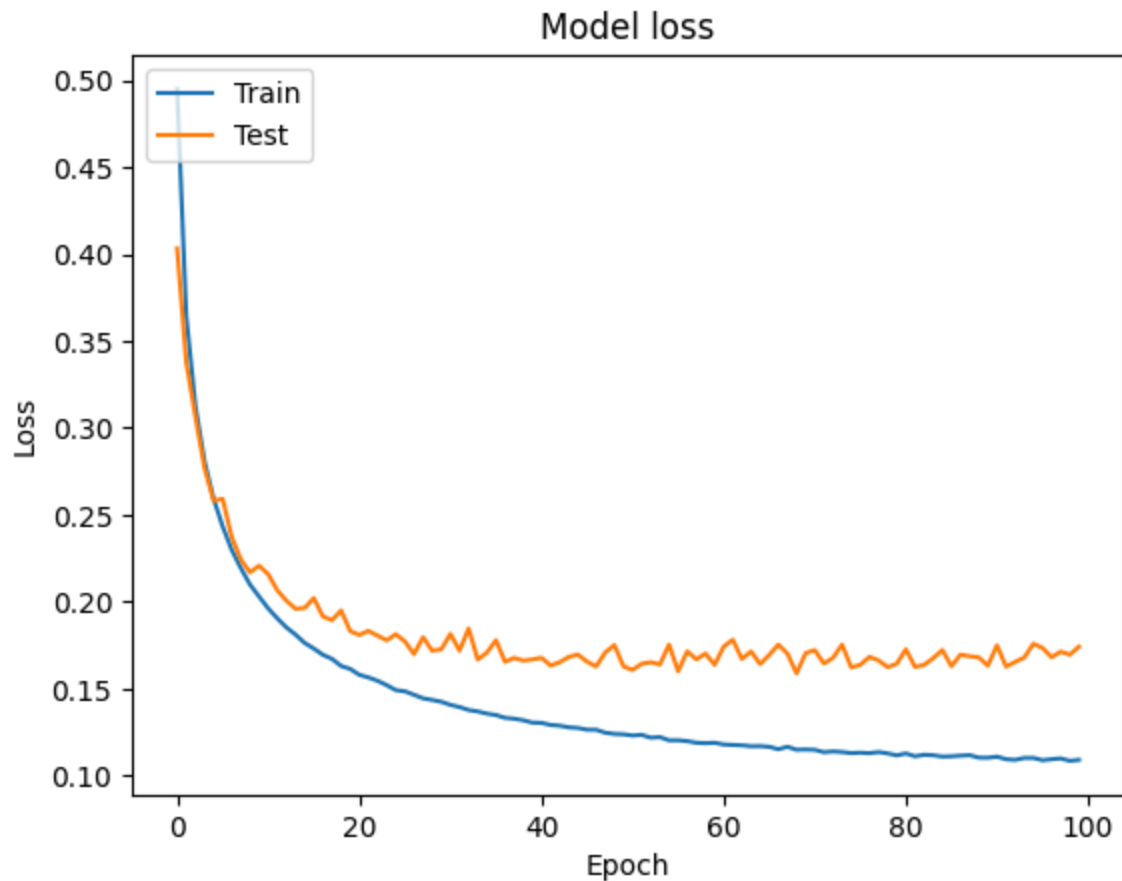
```
In [ ]: import matplotlib.pyplot as plt

# Plot training & validation accuracy values
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
```

```
plt.show()

# Plot training & validation loss values
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()
```





In both graphs we see that the validation results aren't quite as good, which is expected as it is unseen data.

## Step 7: Make Predictions

Making predictions involves feeding new data into the trained model and using the softmax probabilities to determine the most likely class for each instance. This process demonstrates the model's practical utility in classifying new, unseen data.

```
In [ ]: # load the test set
test_set = pd.read_csv('a4-data/test.csv')
test_set.head()
```

```
Out [ ]:
```

	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5	Feature_6	Feature_7	Fea
0	3351	206	27	726	124	3813	192	
1	2732	129	7	212	1	1082	231	
2	2572	24	9	201	25	957	216	
3	2824	69	13	417	39	3223	233	
4	2529	84	5	120	9	1092	227	

5 rows × 54 columns

We need to only scale binary columns like we did in the training data in order to get accurate predictions

```
In [ ]: binary_columns = [col for col in test_set.columns if test_set[col].dropna().count() < 10]
non_binary_columns = [col for col in test_set.columns if col not in binary_columns]

# Normalize non-binary columns
test_set[non_binary_columns] = (test_set[non_binary_columns] - test_set[non_binary_columns].min()) / (test_set[non_binary_columns].max() - test_set[non_binary_columns].min())
test_set.head()
```

```
Out [ ]:
```

	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5	Feature_6	Feature_7	Fea
0	1.397080	0.448163	1.717586	2.142548	1.325494	0.939452	-0.749856	1.
1	-0.817890	-0.239629	-0.950993	-0.271829	-0.781223	-0.813558	0.704324	0
2	-1.390419	-1.177526	-0.684135	-0.323499	-0.370156	-0.893795	0.145024	-0.
3	-0.488686	-0.775570	-0.150419	0.691103	-0.130367	0.560735	0.778898	-0
4	-1.544286	-0.641585	-1.217851	-0.703975	-0.644200	-0.807139	0.555178	0

5 rows × 54 columns

Now we use `model.predict()` to get a matrix with 7 columns and each row representing a probability distribution of the classes, we then get the highest probabilities index which is our prediction

```
In [ ]: import numpy as np

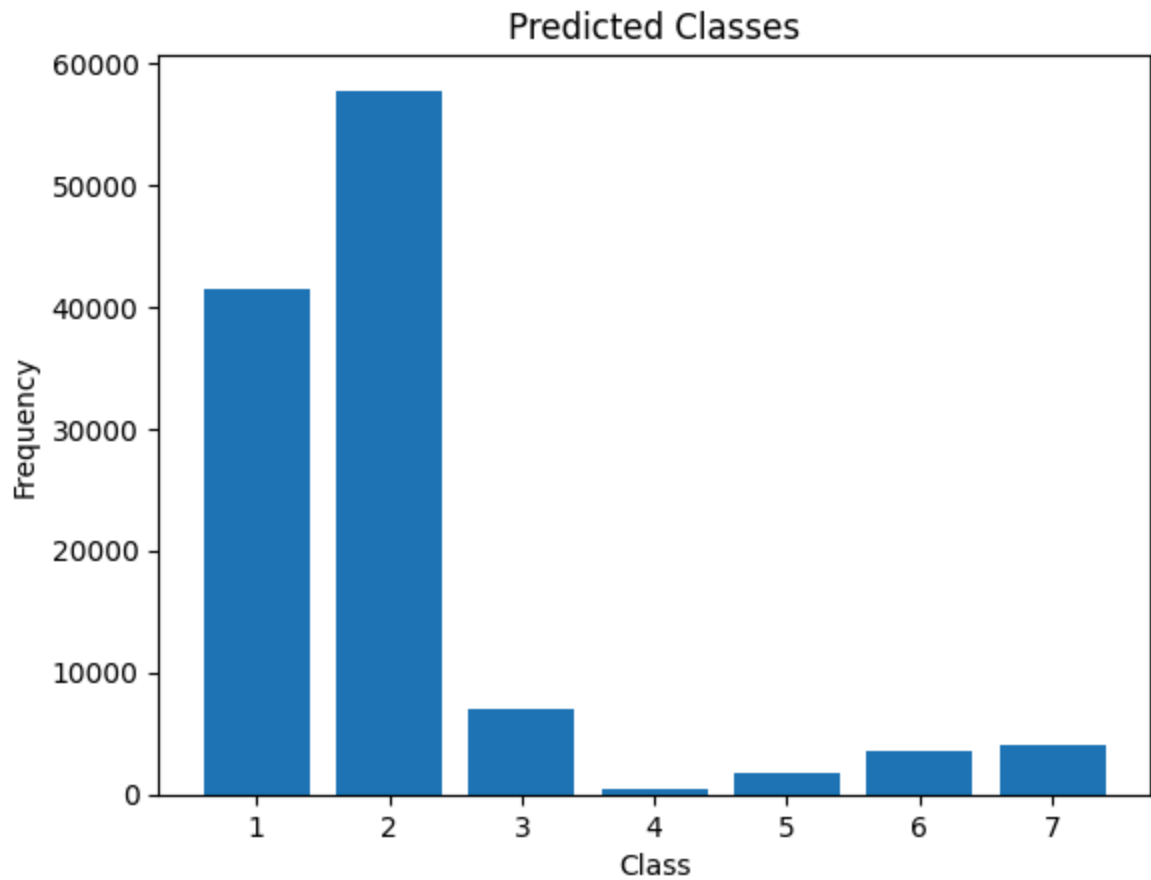
# make predictions and convert them to classes
predictions = model.predict(test_set)
predicted_classes = np.argmax(predictions, axis=1)
```

3632/3632 ————— 1s 293us/step

```
In [ ]: import matplotlib.pyplot as plt

# plot the predicted classes
plt.hist(predicted_classes, bins=range(1, 9), rwidth=0.8, align='left')
```

```
plt.title('Predicted Classes')  
plt.xlabel('Class')  
plt.ylabel('Frequency')  
plt.show()
```



Here we can see the the a large majority of predictions were either 1 or 2, this is a useful insight

## Step 8: Generate Submission File

```
In [ ]: # save the predictions to a CSV file  
submission = pd.DataFrame({'Target': predicted_classes})  
submission.to_csv('submission.csv')
```